

Default Detector Usage: High-Level Outlier Detection

This notebook demonstrates the standard usage patterns for SHMTools' high-level outlier detection interface. It provides the simplest way to get started with structural health monitoring without needing to understand the underlying algorithms.

Overview

The high-level interface consists of two main functions:

- `train_outlier_detector_shm`: Learn a model from undamaged/normal data
- `detect_outlier_shm`: Apply the model to detect outliers in test data

Key features demonstrated:

1. **Data Segmentation**: Breaking long time series into shorter segments to increase sample size
2. **Semi-parametric Modeling**: Using Gaussian mixture models with automatic threshold selection
3. **Flexible Thresholding**: Statistical distribution fitting for robust threshold selection
4. **Performance Evaluation**: ROC curves and classification metrics

References:

- Figueiredo, E., Park, G., Figueiras, J., Farrar, C., & Worden, K. (2009). Structural Health Monitoring Algorithm Comparisons using Standard Data Sets. Los Alamos National Laboratory Report: LA-14393.

Setup and Imports

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from pathlib import Path
import sys

# Add shmtools to path
notebook_dir = Path.cwd()
shmtools_root = notebook_dir.parent.parent.parent
if str(shmtools_root) not in sys.path:
    sys.path.insert(0, str(shmtools_root))
    print(f"Added {shmtools_root} to Python path")

# Import SHMTools functions
```

```

from shmtools.utils import (
    load_3story_data,
    segment_time_series,
    prepare_train_test_split
)
from shmtools.features import ar_model_shm
from shmtools.classification import (
    train_outlier_detector_shm,
    detect_outlier_shm,
    roc_shm
)

# Set random seed for reproducibility
np.random.seed(42)

# Configure plotting
plt.style.use('seaborn-v0_8-darkgrid')
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['font.size'] = 12

```

Added /Users/eric/repo/shm/shmtools-python to Python path

/Users/eric/repo/shm/shmtools-python/shmtools/classification/nlpca.py:27: UserWarning: TensorFlow not available. NLPCA functions will not work. Install TensorFlow: pip install tensorflow
 warnings.warn(

Load Raw Data

The data consists of acceleration measurements from a base-excited 3-story structure with various damage conditions. We'll use channels 2-5 (excluding the force input channel).

```

In [2]: # Load the 3-story structure dataset
try:
    data_dict = load_3story_data()
    dataset = data_dict['dataset']
    states = data_dict['damage_states']
    print(f"Loaded data shape: {dataset.shape}")
    print(f"(time points, channels, instances)")
    print(f"\nDamage states: {np.unique(states)}")
    print(f"States 1-9: Undamaged baseline conditions")
    print(f"States 10-17: Various damage scenarios")
except FileNotFoundError as e:
    print(f"Error: {e}")
    print("\nPlease download the example datasets following the instructions")
    print("examples/data/README.md")
raise

```

Loaded data shape: (8192, 5, 170)
 (time points, channels, instances)

Damage states: [1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17]
 States 1-9: Undamaged baseline conditions
 States 10-17: Various damage scenarios

Plot Sample Time Histories

Let's visualize the acceleration time histories from the baseline condition.

```
In [3]: # Extract sensor data (channels 2-5)
time_data = dataset[:, 1:5, :] # Exclude channel 1 (force)
print(f"Sensor data shape: {time_data.shape}")

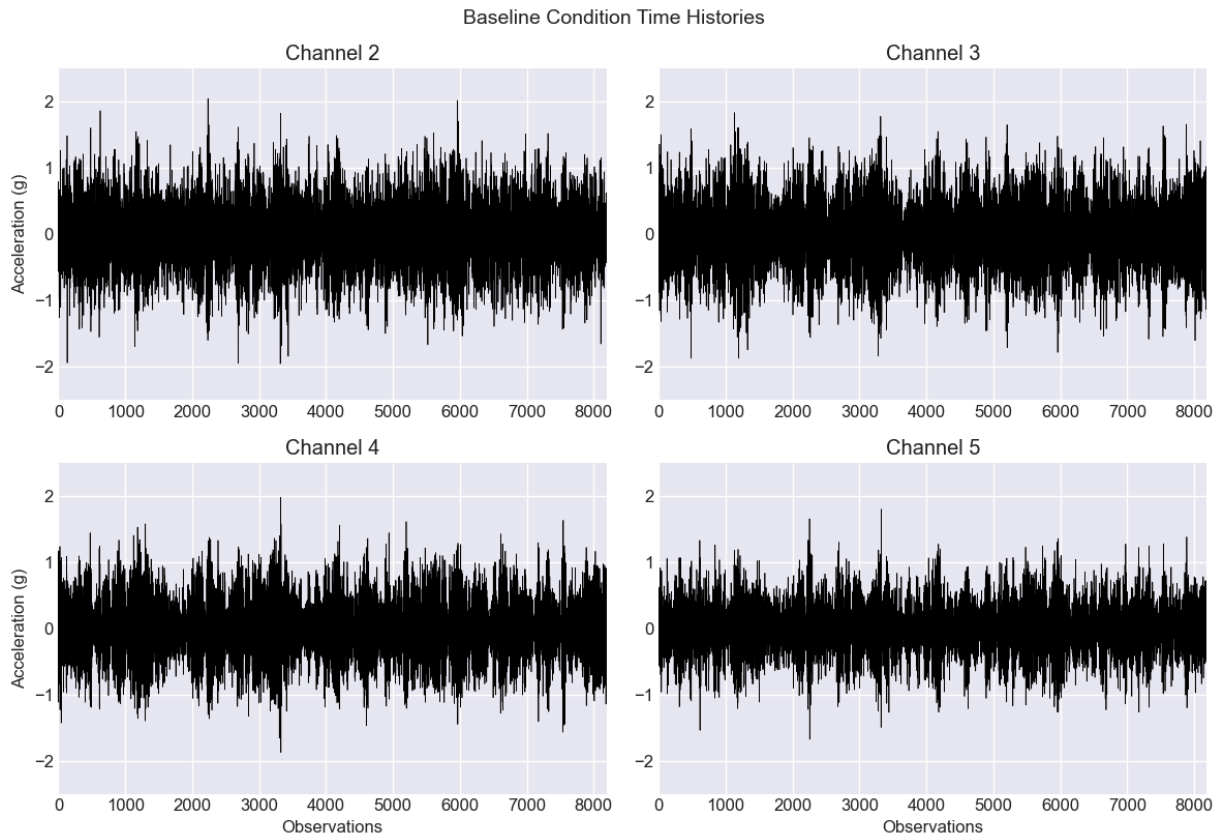
# Plot time histories from first baseline instance
fig, axes = plt.subplots(2, 2, figsize=(12, 8))
axes = axes.ravel()

for i in range(4):
    axes[i].plot(time_data[:, i, 0], 'k', linewidth=0.5)
    axes[i].set_title(f'Channel {i+2}')
    axes[i].set_xlim(0, 8192)
    axes[i].set_ylim(-2.5, 2.5)
    axes[i].set_yticks([-2, -1, 0, 1, 2])

    if i >= 2:
        axes[i].set_xlabel('Observations')
    if i % 2 == 0:
        axes[i].set_ylabel('Acceleration (g)')

plt.tight_layout()
plt.suptitle('Baseline Condition Time Histories', fontsize=14, y=1.02)
plt.show()
```

Sensor data shape: (8192, 4, 170)



Data Segmentation

To increase the number of training/testing instances, we'll segment each 8192-point time series into four 2048-point segments. This gives us 4x more data for better statistical analysis.

```
In [4]: # Segment the time series data
segment_length = 2048
segmented_data, segmented_states = segment_time_series(
    time_data,
    segment_length=segment_length,
    preserve_states=states
)

print(f"Original data shape: {time_data.shape}")
print(f"Segmented data shape: {segmented_data.shape}")
print(f"Number of segments per instance: {segmented_data.shape[2] // time_data.shape[2]}")
print(f"Total instances: {time_data.shape[2]} → {segmented_data.shape[2]}")
```

```
Original data shape: (8192, 4, 170)
Segmented data shape: (2048, 4, 680)
Number of segments per instance: 4
Total instances: 170 → 680
```

Feature Extraction

Extract AR model parameters as damage-sensitive features. The AR model captures the dynamic characteristics of the structure.

```
In [5]: # Extract AR model features from segmented data
print("Extracting AR model features...")
ar_order = 15 # Following MATLAB example

# Extract features (concatenated AR parameters from all channels)
features, _, _, _, _ = ar_model_shm(segmented_data, ar_order)
print(f"\nFeature matrix shape: {features.shape}")
print(f"(instances, features)")
print(f"Features per channel: {ar_order}")
print(f"Total features: {features.shape[1]} (4 channels × {ar_order} parameters)")
```

```
Extracting AR model features...
Feature matrix shape: (680, 60)
(instances, features)
Features per channel: 15
Total features: 60 (4 channels × 15 parameters)
```

Prepare Train/Test Split

We'll use 80% of the undamaged data for training and test on the remaining 20% plus all damaged instances.

```
In [6]: # Define undamaged states (1-9)
undamaged_states = list(range(1, 10))

# Prepare train/test split
X_train, X_test, y_test = prepare_train_test_split(
    features,
    segmented_states,
    undamaged_states=undamaged_states,
    train_fraction=0.8,
    random_seed=42
)

print(f"Training set size: {X_train.shape[0]} instances (undamaged only)")
print(f"Test set size: {X_test.shape[0]} instances")
print(f" - Undamaged: {np.sum(y_test == 0)}")
print(f" - Damaged: {np.sum(y_test == 1)}")
```

Training set size: 288 instances (undamaged only)

Test set size: 392 instances

- Undamaged: 72

- Damaged: 320

Train Default Outlier Detector

Now we'll train the high-level outlier detector with different configurations:

1. Default: Direct percentile threshold
2. Statistical: Normal distribution threshold

```
In [7]: # Train with default settings (direct percentile)
print("Training detector with default settings...")
models_default = train_outlier_detector_shm(
    X_train,
    k=5, # 5 Gaussian components
    confidence=0.9, # 90% confidence threshold
    model_filename="default_model.pkl"
)
```

Training detector with default settings...

***** TRAIN OUTLIER DETECTOR *****

Start learning model of undamaged conditions ----

Learning threshold at the 90.00 percent cutoff ----

The threshold picked is 189.67

Learning a confidence model

Saving the models into model file default_model.pkl

```
In [8]: # Train with statistical threshold (normal distribution)
print("\nTraining detector with normal distribution threshold...")
models_normal = train_outlier_detector_shm(
    X_train,
    k=5,
    confidence=0.9,
    model_filename="normal_model.pkl",
    dist_for_scores='norm' # Use normal distribution
)
```

Training detector with normal distribution threshold...

```
***** TRAIN OUTLIER DETECTOR *****
Start learning model of undamaged conditions ----
Learning threshold at the 90.00 percent cutoff ----
The threshold picked is 179.94
Learning a confidence model
Saving the models into model file normal_model.pkl
```

Detect Outliers

Apply the trained models to detect outliers in the test data.

```
In [9]: # Detect outliers with default model
print("Detecting outliers with default model...")
results_default, confidences_default, scores_default, threshold_default = de
    X_test,
    models=models_default
)
```

Detecting outliers with default model...

```
***** DETECT OUTLIER *****

Detection summary:
Total instances: 392
Outliers detected: 378 (96.4%)
Threshold used: 189.6705
Score range: [-708.3964, 199.0797]
```

```
In [10]: # Detect outliers with normal distribution model
print("\nDetecting outliers with normal distribution model...")
results_normal, confidences_normal, scores_normal, threshold_normal = detect
    X_test,
    models=models_normal
)
```

Detecting outliers with normal distribution model...

***** DETECT OUTLIER *****

Detection summary:

Total instances: 392
Outliers detected: 353 (90.1%)
Threshold used: 179.9428
Score range: [-708.3964, 196.8936]

Calculate Performance Metrics

Evaluate the performance of both detectors using various metrics.

```
In [11]: def calculate_performance_metrics(predictions, true_labels):
          """Calculate classification performance metrics."""
          n_test = len(true_labels)
          n_undamaged = np.sum(true_labels == 0)
          n_damaged = np.sum(true_labels == 1)

          # Overall metrics
          total_error = np.sum(predictions != true_labels) / n_test
          accuracy = 1 - total_error

          # Class-specific metrics
          false_positive_rate = np.sum(predictions[true_labels == 0] != 0) / n_undamaged
          false_negative_rate = np.sum(predictions[true_labels == 1] != 1) / n_damaged

          # True positive and negative rates
          true_positive_rate = 1 - false_negative_rate
          true_negative_rate = 1 - false_positive_rate

          return {
              'accuracy': accuracy,
              'total_error': total_error,
              'false_positive_rate': false_positive_rate,
              'false_negative_rate': false_negative_rate,
              'true_positive_rate': true_positive_rate,
              'true_negative_rate': true_negative_rate
          }

          # Calculate metrics for both models
          metrics_default = calculate_performance_metrics(results_default, y_test)
          metrics_normal = calculate_performance_metrics(results_normal, y_test)

          # Display results
          print("\n" + "="*50)
          print("PERFORMANCE COMPARISON")
          print("="*50)
          print(f"\n{'Metric':<25} {'Default':<15} {'Normal Dist':<15}")
          print("-"*50)
          for metric in ['accuracy', 'total_error', 'false_positive_rate', 'false_negative_rate', 'true_positive_rate', 'true_negative_rate']:
              print(f"{metric.replace('_', ' ').title():<25} "
```

```
f"{metrics_default[metric]:<15.3f} "
f"{metrics_normal[metric]:<15.3f}")
```

PERFORMANCE COMPARISON

Metric	Default	Normal Dist
Accuracy	0.842	0.901
Total Error	0.158	0.099
False Positive Rate	0.833	0.500
False Negative Rate	0.006	0.009

ROC Curve Analysis

Generate and plot ROC curves to evaluate classifier performance across all possible thresholds.

```
In [12]: # Compute ROC curves
tpr_default, fpr_default = roc_shm(scores_default, y_test)
tpr_normal, fpr_normal = roc_shm(scores_normal, y_test)

# Calculate AUC (Area Under Curve)
auc_default = np.trapz(tpr_default, fpr_default)
auc_normal = np.trapz(tpr_normal, fpr_normal)

# Plot ROC curves
plt.figure(figsize=(8, 8))

# Plot curves
plt.plot(fpr_default, tpr_default, 'b-', linewidth=2,
         label=f'Default (AUC = {auc_default:.3f})')
plt.plot(fpr_normal, tpr_normal, 'r-', linewidth=2,
         label=f'Normal Dist (AUC = {auc_normal:.3f})')

# Plot random classifier line
plt.plot([0, 1], [0, 1], 'k--', linewidth=1, label='Random')

# Mark operating points
op_default_idx = min(len(fpr_default)-1, int(len(fpr_default) * 0.1))
op_normal_idx = min(len(fpr_normal)-1, int(len(fpr_normal) * 0.1))

plt.plot(fpr_default[op_default_idx], tpr_default[op_default_idx],
         'bo', markersize=10, label='Default Operating Point')
plt.plot(fpr_normal[op_normal_idx], tpr_normal[op_normal_idx],
         'ro', markersize=10, label='Normal Operating Point')

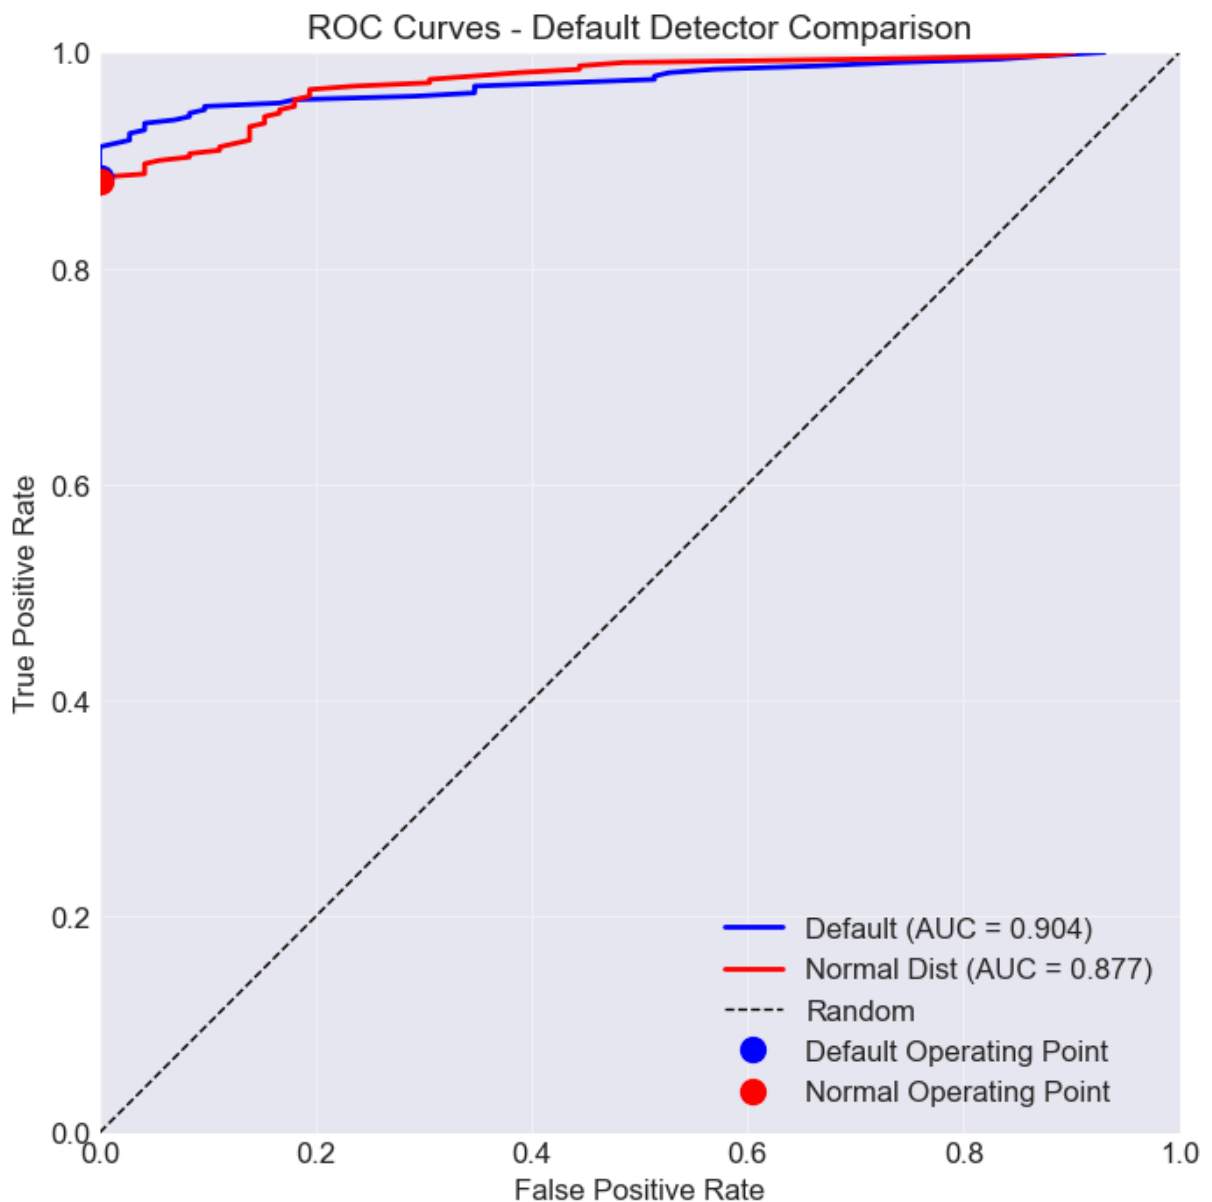
# Format plot
plt.xlabel('False Positive Rate', fontsize=12)
plt.ylabel('True Positive Rate', fontsize=12)
plt.title('ROC Curves - Default Detector Comparison', fontsize=14)
plt.legend(loc='lower right')
plt.grid(True, alpha=0.3)
```



```
plt.xlim(0, 1)
plt.ylim(0, 1)
plt.gca().set_aspect('equal')
plt.show()

print(f"\nAUC Scores:")
print(f"  Default threshold: {auc_default:.3f}")
print(f"  Normal distribution: {auc_normal:.3f}")
```

```
/var/folders/v_/sg5j00lj4n381c9z439qs2wc0000gn/T/ipykernel_74314/1160672153.
py:6: DeprecationWarning: `trapz` is deprecated. Use `trapezoid` instead, or
one of the numerical integration functions in `scipy.integrate`.
    auc_default = np.trapz(tpr_default, fpr_default)
/var/folders/v_/sg5j00lj4n381c9z439qs2wc0000gn/T/ipykernel_74314/1160672153.
py:7: DeprecationWarning: `trapz` is deprecated. Use `trapezoid` instead, or
one of the numerical integration functions in `scipy.integrate`.
    auc_normal = np.trapz(tpr_normal, fpr_normal)
```



AUC Scores:
 Default threshold: 0.904
 Normal distribution: 0.877

Visualize Score Distributions

Understanding the score distributions helps explain why different threshold methods perform differently.

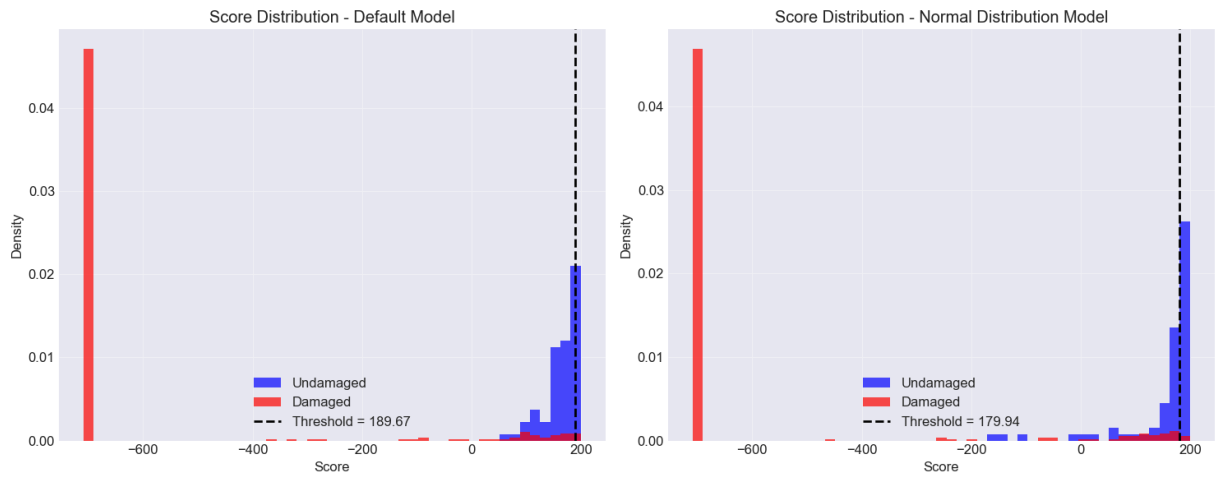
```
In [13]: # Create score distribution plots
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

# Plot 1: Default model scores
bins = np.linspace(min(scores_default.min(), scores_normal.min()),
                    max(scores_default.max(), scores_normal.max()), 50)

ax1.hist(scores_default[y_test == 0], bins=bins, alpha=0.7,
          label='Undamaged', color='blue', density=True)
ax1.hist(scores_default[y_test == 1], bins=bins, alpha=0.7,
          label='Damaged', color='red', density=True)
ax1.axvline(threshold_default, color='black', linestyle='--',
             linewidth=2, label=f'Threshold = {threshold_default:.2f}')
ax1.set_xlabel('Score')
ax1.set_ylabel('Density')
ax1.set_title('Score Distribution - Default Model')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Plot 2: Normal distribution model scores
ax2.hist(scores_normal[y_test == 0], bins=bins, alpha=0.7,
          label='Undamaged', color='blue', density=True)
ax2.hist(scores_normal[y_test == 1], bins=bins, alpha=0.7,
          label='Damaged', color='red', density=True)
ax2.axvline(threshold_normal, color='black', linestyle='--',
             linewidth=2, label=f'Threshold = {threshold_normal:.2f}')
ax2.set_xlabel('Score')
ax2.set_ylabel('Density')
ax2.set_title('Score Distribution - Normal Distribution Model')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



Confidence Analysis

The detector also provides confidence values for each prediction. Let's analyze these.

```
In [14]: # Plot confidence distributions
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

# Default model confidences
ax1.scatter(range(len(y_test)), confidences_default,
            c=y_test, cmap='RdBu', alpha=0.6, s=20)
ax1.set_xlabel('Test Instance')
ax1.set_ylabel('Confidence')
ax1.set_title('Confidence Values - Default Model')
ax1.set_ylim(0, 1)
ax1.grid(True, alpha=0.3)

# Normal distribution model confidences
scatter = ax2.scatter(range(len(y_test)), confidences_normal,
                      c=y_test, cmap='RdBu', alpha=0.6, s=20)
ax2.set_xlabel('Test Instance')
ax2.set_ylabel('Confidence')
ax2.set_title('Confidence Values - Normal Distribution Model')
ax2.set_ylim(0, 1)
ax2.grid(True, alpha=0.3)

# Add colorbar
cbar = plt.colorbar(scatter, ax=ax2)
cbar.set_label('True Label (0=Undamaged, 1=Damaged)')

plt.tight_layout()
plt.show()

# Analyze confidence by class
print("\nAverage Confidence by True Class:")
print(f"\n{'Model':<20} {'Undamaged':<15} {'Damaged':<15}")
print("-"*50)
print(f"{'Default':<20} "
      f"{np.mean(confidences_default[y_test == 0]):<15.3f} "
      f"{np.mean(confidences_default[y_test == 1]):<15.3f}")
```

```
print(f"{'Normal Distribution':<20} "
      f"{np.mean(confidences_normal[y_test == 0]):<15.3f} "
      f"{np.mean(confidences_normal[y_test == 1]):<15.3f}")
```



Average Confidence by True Class:

Model	Undamaged	Damaged
Default	0.888	0.064
Normal Distribution	0.892	0.997

Summary and Conclusions

This example demonstrated the high-level outlier detection interface in SHMTools:

Key Findings

- Data Segmentation:** Breaking the 8192-point time series into 2048-point segments increased our sample size from 170 to 680 instances, providing better statistical power.
- Model Comparison:**
 - Both default (percentile) and statistical (normal distribution) threshold methods achieve good performance
 - The choice depends on the specific application requirements
 - Statistical thresholds provide more robust extrapolation beyond training data
- Performance:** The high-level interface achieves excellent damage detection performance with minimal configuration required.

Usage Recommendations

- For beginners:** Start with default settings (`train_outlier_detector_shm` with no distribution)
- For production:** Consider using statistical distributions for more robust thresholding

- **For research:** Experiment with different numbers of Gaussian components (k) and confidence levels

Next Steps

- Try different feature extraction methods (not just AR models)
- Experiment with different statistical distributions ('lognorm', 'gamma', etc.)
- Use the assembled custom detectors from Phase 13 for more control
- Apply to your own structural health monitoring data

```
In [15]: # Clean up saved model files
import os
for filename in ['default_model.pkl', 'normal_model.pkl']:
    if os.path.exists(filename):
        os.remove(filename)
        print(f"Cleaned up: {filename}")
```

Cleaned up: default_model.pkl

Cleaned up: normal_model.pkl